


Assignment-3


Name: K Naga Sai Krishna

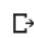
```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
from google.colab import files
uploaded = files.upload()
```

 penguins_size.csv
penguins_size.csv(application/vnd.ms-excel) - 13519 bytes, last modified: n/a - 100% done
Saving penguins_size.csv to penguins_size.csv

```
import io
df = pd.read_csv(io.BytesIO(uploaded['penguins_size.csv']))
```

 `df.head()`



	species	island	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	sex
0	Adelie	Torgersen	39.1	18.7	181.0	3750.0	MALE
1	Adelie	Torgersen	39.5	17.4	186.0	3800.0	FEMALE
2	Adelie	Torgersen	40.3	18.0	195.0	3250.0	FEMALE
3	Adelie	Torgersen	NaN	NaN	NaN	NaN	NaN
4	Adelie	Torgersen	36.7	19.3	193.0	3450.0	FEMALE

```
[7] df.shape
```

```
(344, 7)
```

```
[8] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 344 entries, 0 to 343
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   species                344 non-null    object
1   island                 344 non-null    object
2   culmen_length_mm       342 non-null    float64
3   culmen_depth_mm        342 non-null    float64
4   flipper_length_mm      342 non-null    float64
5   body_mass_g            342 non-null    float64
6   sex                    334 non-null    object
dtypes: float64(4), object(3)
memory usage: 18.9+ KB
```

Dropping the columns that have NaN in 1 to 5 Columns


```
df.dropna(subset=df.columns[:6],inplace=True)
```

```
[10] df.shape
```

```
(342, 7)
```

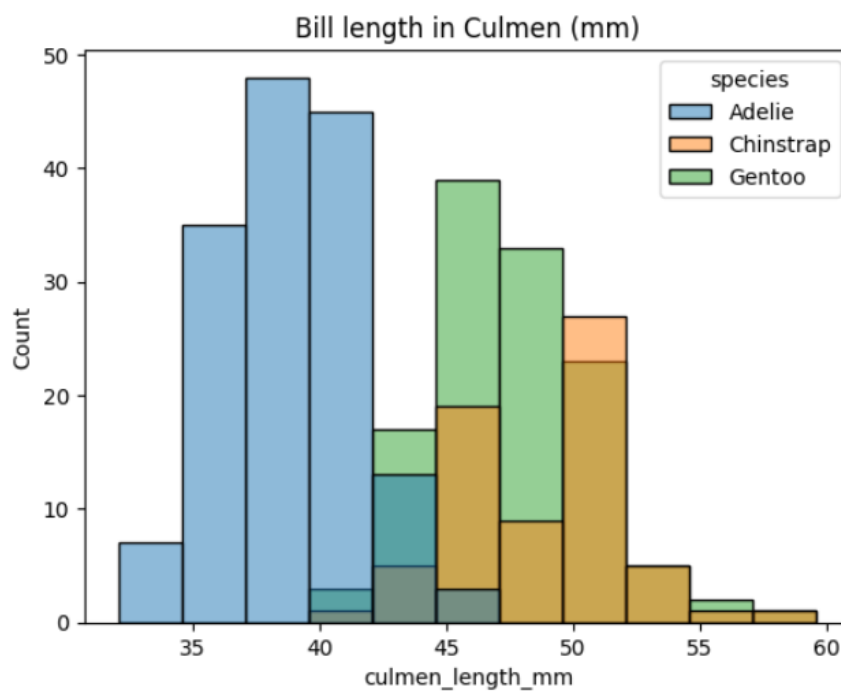
Univariate Analysis

```
[18] sns.histplot(df,x='flipper_length_mm',hue='species').set(title='Flipper Length in all species');
```



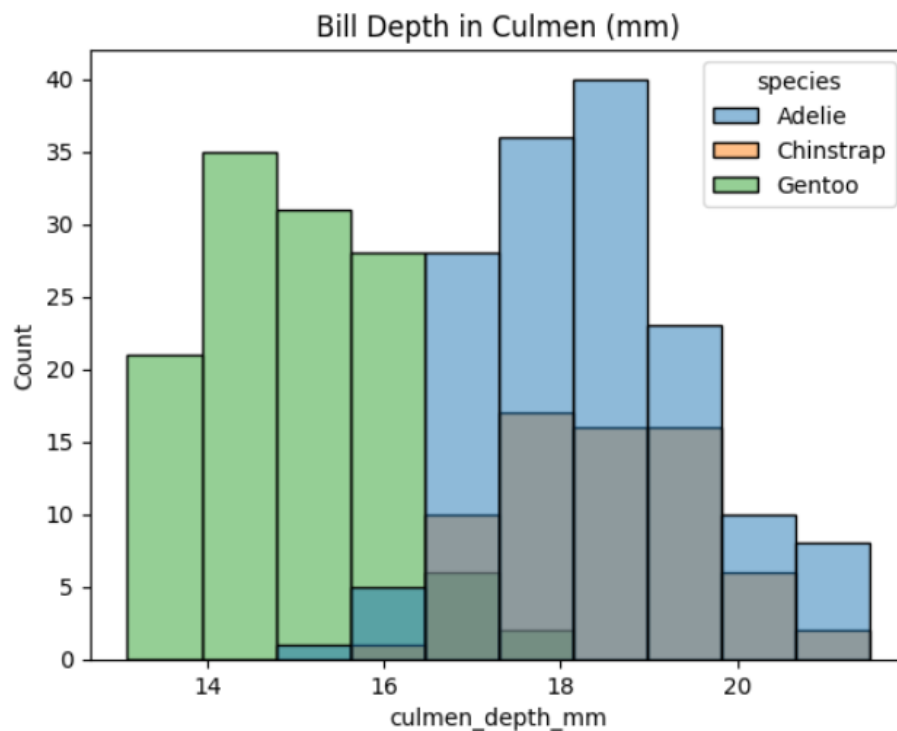
```
sns.histplot(df,x='culmen_length_mm',hue='species').set(title='Bill length in Culmen (mm)')
```

```
[Text(0.5, 1.0, 'Bill length in Culmen (mm)')]
```




```
[20] sns.histplot(df,x='culmen_depth_mm',hue='species').set(title='Bill Depth in Culmen (mm)')
```

```
[Text(0.5, 1.0, 'Bill Depth in Culmen (mm)')]
```



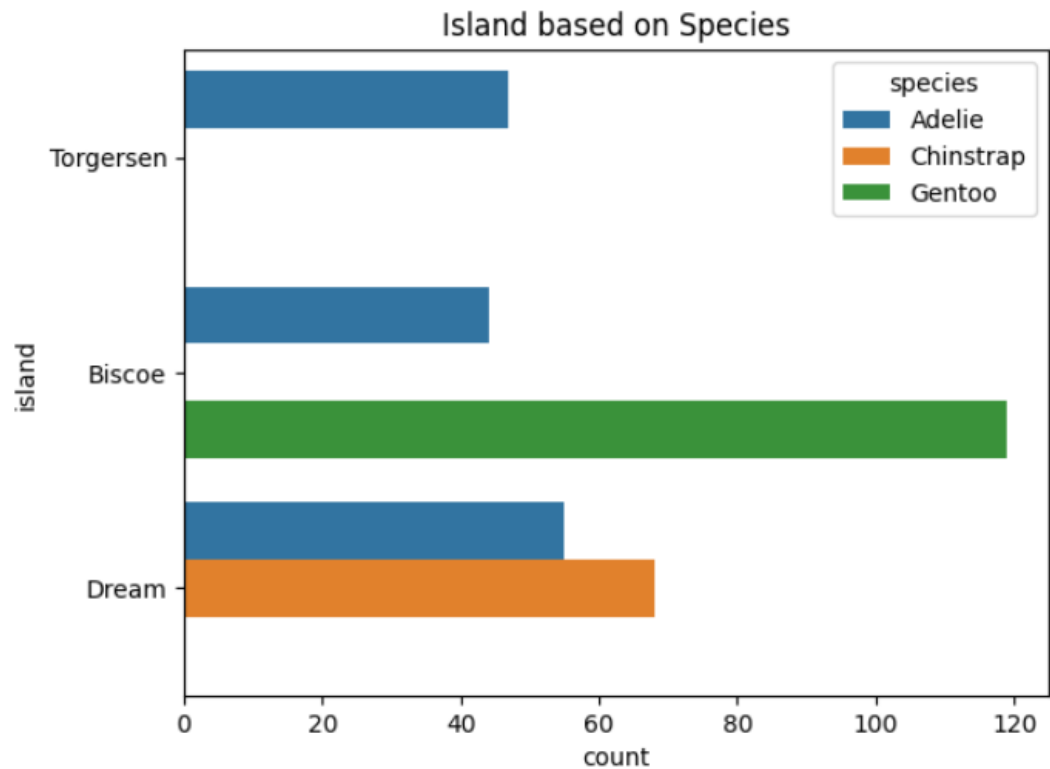
```
[21] sns.histplot(df,x='body_mass_g',hue='species').set(title='Body Mass in all species')
```

```
[Text(0.5, 1.0, 'Body Mass in all species')]
```



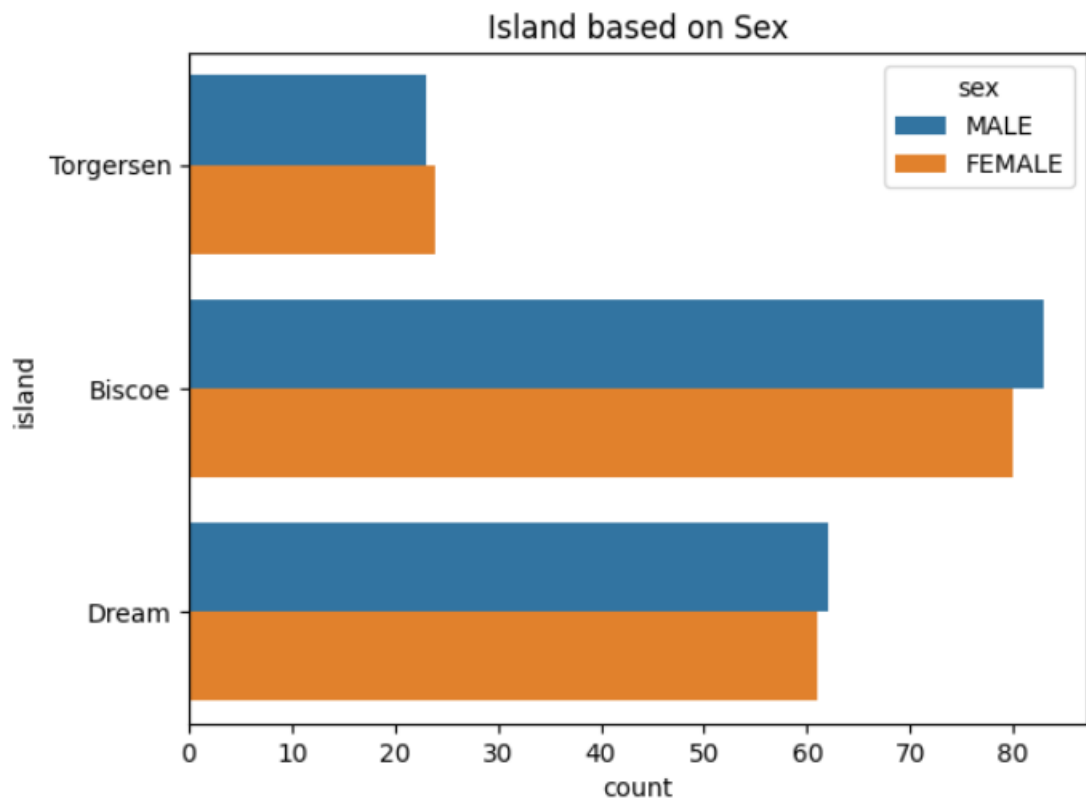
```
[25] sns.countplot(y='island',data=df,hue='species').set(title='Island based on Species')
```

```
[Text(0.5, 1.0, 'Island based on Species')]
```



```
▶ sns.countplot(y='island',data=df,hue='sex').set(title='Island based on Sex')
```


```
↳ [Text(0.5, 1.0, 'Island based on Sex')]
```



Bivariate Analysis


```
[76] sns.scatterplot(data=df,x='culmen_length_mm',y='culmen_depth_mm',hue='species')
```

<Axes: xlabel='culmen_length_mm', ylabel='culmen_depth_mm'>




```
sns.boxplot(data=df, x='species', y='culmen_depth_mm')
```

<Axes: xlabel='species', ylabel='culmen_depth_mm'>



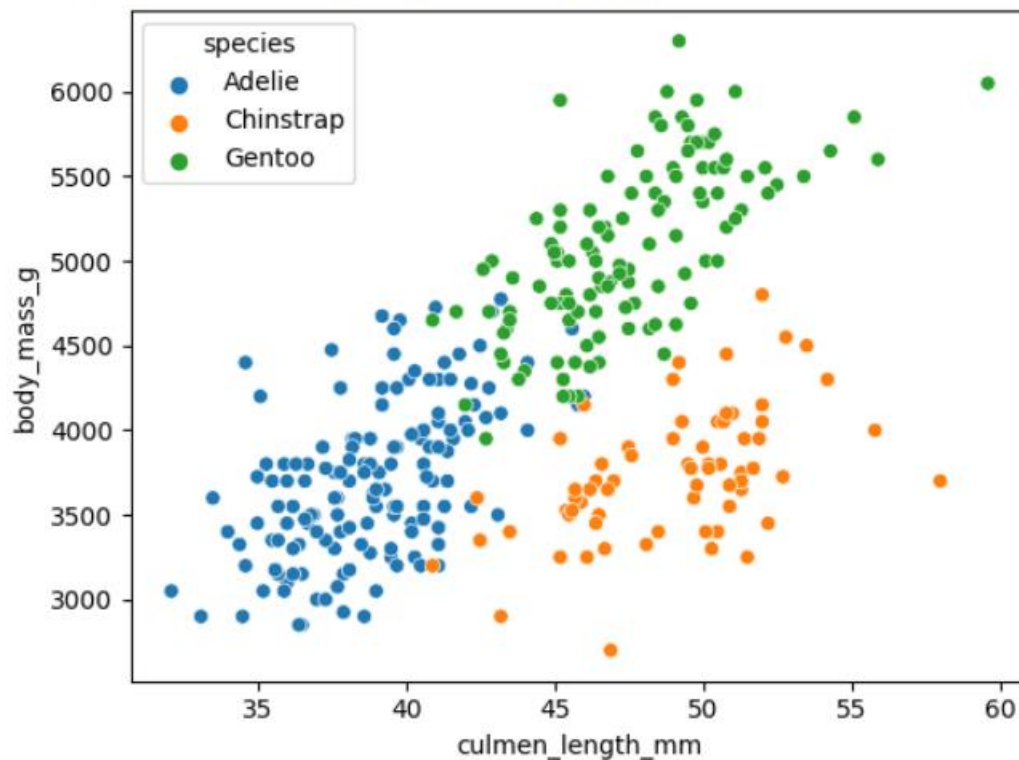
```
sns.boxplot(data=df, x='species', y='culmen_length_mm')
```

```
<Axes: xlabel='species', ylabel='culmen_length_mm'>
```



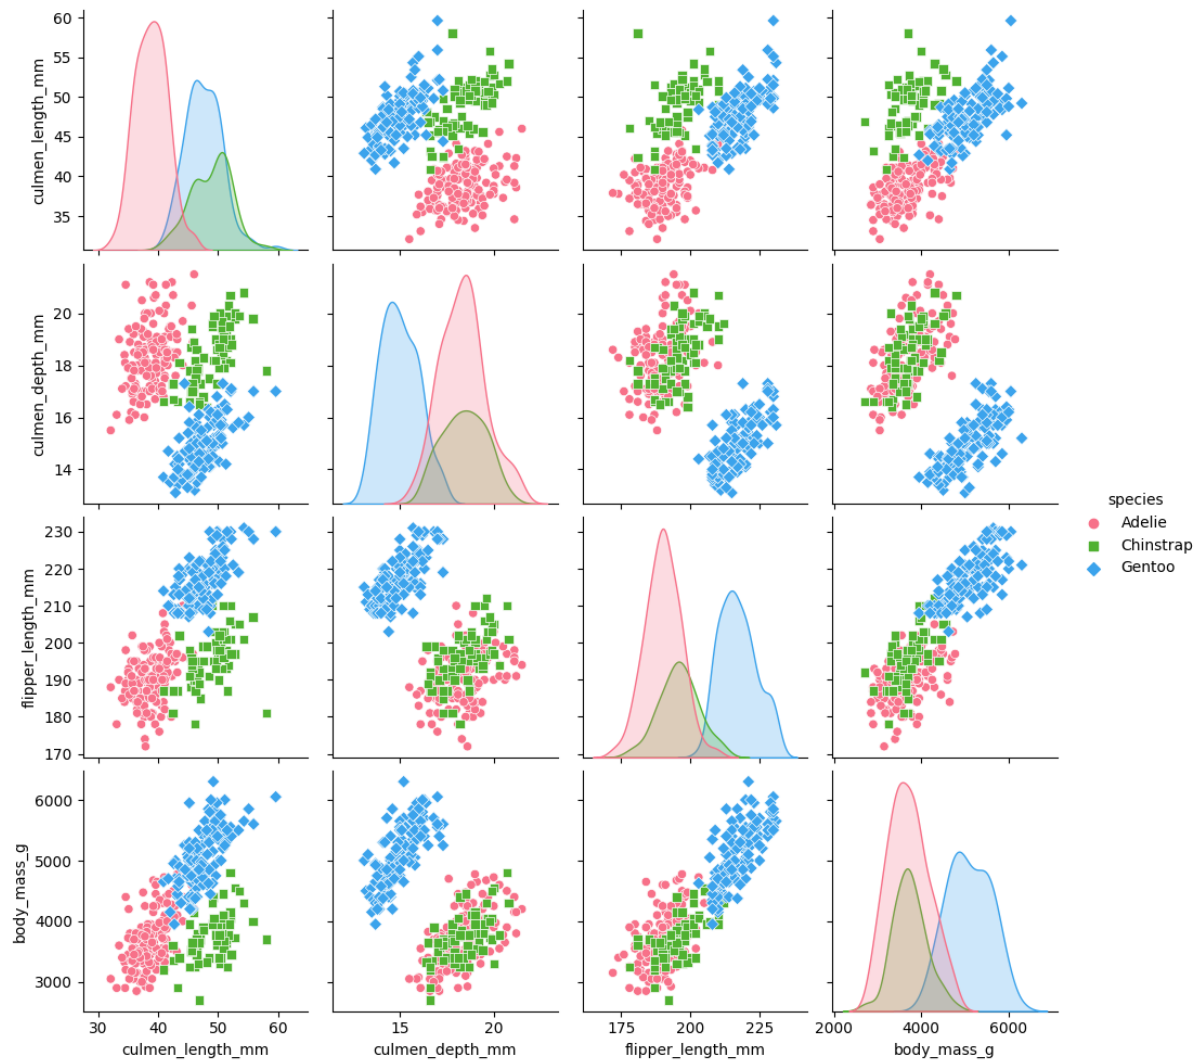
```
[35] sns.scatterplot(data=df, x='culmen_length_mm', y='body_mass_g', hue='species')
```

```
<Axes: xlabel='culmen_length_mm', ylabel='body_mass_g'>
```



Multivariate Analysis

```
sns.pairplot(df, hue='species', palette='husl', markers=['o', 's', 'D']);
```



Descriptive Statistics

```
[30] df.describe()
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g
count	333.000000	333.000000	333.000000	333.000000
mean	43.992793	17.164865	200.966967	4207.057057
std	5.468668	1.969235	14.015765	805.215802
min	32.100000	13.100000	172.000000	2700.000000
25%	39.500000	15.600000	190.000000	3550.000000
50%	44.500000	17.300000	197.000000	4050.000000
75%	48.600000	18.700000	213.000000	4775.000000
max	59.600000	21.500000	231.000000	6300.000000

```
[14] df['species'].value_counts()
```

```
Adelie      151
Gentoo      123
Chinstrap    68
Name: species, dtype: int64
```

```
adelie_sp = df[df['species']=='Adelie']
adelie_sp.describe()
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g
count	151.000000	151.000000	151.000000	151.000000
mean	38.791391	18.346358	189.953642	3700.662252
std	2.663405	1.216650	6.539457	458.566126
min	32.100000	15.500000	172.000000	2850.000000
25%	36.750000	17.500000	186.000000	3350.000000
50%	38.800000	18.400000	190.000000	3700.000000
75%	40.750000	19.000000	195.000000	4000.000000
max	46.000000	21.500000	210.000000	4775.000000


```
gentoo_sp = df[df['species']=='Gentoo']
gentoo_sp.describe()
```



	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g
count	123.000000	123.000000	123.000000	123.000000
mean	47.504878	14.982114	217.186992	5076.016260
std	3.081857	0.981220	6.484976	504.116237
min	40.900000	13.100000	203.000000	3950.000000
25%	45.300000	14.200000	212.000000	4700.000000
50%	47.300000	15.000000	216.000000	5000.000000
75%	49.550000	15.700000	221.000000	5500.000000
max	59.600000	17.300000	231.000000	6300.000000

```
[17] chinstrap_sp = df[df['species']=='Chinstrap']
chinstrap_sp.describe()
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g
count	68.000000	68.000000	68.000000	68.000000
mean	48.833824	18.420588	195.823529	3733.088235
std	3.339256	1.135395	7.131894	384.335081
min	40.900000	16.400000	178.000000	2700.000000
25%	46.350000	17.500000	191.000000	3487.500000
50%	49.550000	18.450000	196.000000	3700.000000
75%	51.075000	19.400000	201.000000	3950.000000
max	58.000000	20.800000	212.000000	4800.000000

Handling missing values

```
[22] df['sex'].unique()
```

```
array(['MALE', 'FEMALE', nan, '.'], dtype=object)
```

```
[23] df.loc[df['sex']=='.', ['sex']] = np.nan
```

```
df.dropna(subset=['sex'], inplace=True)
```

```
[24] df['sex'].unique()
```

```
array(['MALE', 'FEMALE'], dtype=object)
```


```
[68] df.isnull().sum()
```

```
species      0
island        0
culmen_length_mm  0
culmen_depth_mm  0
flipper_length_mm  0
body_mass_g   0
sex           0
dtype: int64
```

Checking for Outliers by plotting


```
[41] sns.boxplot(df.culmen_depth_mm)
```

<Axes: >




```
▶ sns.boxplot(df.culmen_length_mm)
```

↗ <Axes: >




```
[43] sns.boxplot(df.flipper_length_mm)
```

<Axes: >



```
[44] sns.boxplot(df.body_mass_g)
```

<Axes: >



No outliers were found

Corelation with target column

```
[51] df.corr()
```

```
<ipython-input-51-2f6f6606aa2c>:1: FutureWarning: The default value of numeric_only in  
df.corr()
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g
culmen_length_mm	1.000000	-0.228626	0.653096	0.589451
culmen_depth_mm	-0.228626	1.000000	-0.577792	-0.472016
flipper_length_mm	0.653096	-0.577792	1.000000	0.872979
body_mass_g	0.589451	-0.472016	0.872979	1.000000

```
df.corr()['species']
```

```
species      1.000000  
island      -0.622428  
culmen_length_mm  0.730548  
culmen_depth_mm -0.740346  
flipper_length_mm  0.850737  
body_mass_g   0.750434  
sex          0.010964  
Name: species, dtype: float64
```

Label Encoding

```
[53] from sklearn.preprocessing import LabelEncoder  
le = LabelEncoder()
```

```
[54] df.sex = le.fit_transform(df.sex)  
df.species = le.fit_transform(df.species)  
df.island = le.fit_transform(df.island)
```

```
[55] df.head()
```

	species	island	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	sex
0	0	2	39.1	18.7	181.0	3750.0	1
1	0	2	39.5	17.4	186.0	3800.0	0
2	0	2	40.3	18.0	195.0	3250.0	0
4	0	2	36.7	19.3	193.0	3450.0	0
5	0	2	39.3	20.6	190.0	3650.0	1

Splitting

```
[60] X = df.drop('species', axis=1)
      y = df['species']
```

Scaling

```
[56] from sklearn.preprocessing import MinMaxScaler
      scale =MinMaxScaler()
```

```
[59] X_scaled= pd.DataFrame(scale.fit_transform(X),columns =X.columns)
      X_scaled.head()
```

	island	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	sex
0	1.0	0.254545	0.666667	0.152542	0.291667	1.0
1	1.0	0.269091	0.511905	0.237288	0.305556	0.0
2	1.0	0.298182	0.583333	0.389831	0.152778	0.0
3	1.0	0.167273	0.738095	0.355932	0.208333	0.0
4	1.0	0.261818	0.892857	0.305085	0.263889	1.0

S

Train Test Split and Checking the sizes

```
[66] from sklearn.model_selection import train_test_split
      X_train,X_test,y_train,y_test = train_test_split(X_scaled,y,test_size=0.2,random_state=42)
```

```
[71] print("X_train shape=",X_train.shape,"\nX_test shape=",X_test.shape,
        "\ny_train shape=",y_train.shape,"\ny_test shape=",y_test.shape)
```

```
X_train shape= (266, 6)
X_test shape= (67, 6)
y_train shape= (266,)
y_test shape= (67,)
```